

Semester	B.E. Semester VII
Subject	Deep Learning
Subject Professor In- charge	Dr. Nayana Mahajan
Laboratory	M201B

Student Name	Harsh Jain	Division	B
Roll Number	22108B0054	Batch	4
Grade and Subject Teacher's Signature			

Experiment Number	2
Experiment Title	<p>To train and evaluate a single-layer feedforward neural network on a real-world binary classification dataset using Stochastic Gradient Descent (SGD) and Momentum-based Gradient Descent (Momentum GD) as optimization techniques. The objective is to compare and analyze the performance of both optimizers in terms of:</p> <ul style="list-style-type: none"> Convergence Rate: How quickly the training loss decreases over epochs. Training Speed: The computational efficiency and time taken during training. Classification Accuracy: The predictive performance on unseen test data. <p>This study aims to highlight the impact of optimization strategy on neural network training effectiveness, particularly in low-complexity models such as single-layer networks.</p>
Resources / Apparatus Required	Software: Google Colab
Algorithm	<ol style="list-style-type: none"> Load Dataset: Create Binary Target: One-Hot Encode Categorical Features: Prepare Features and Labels: Normalize Features: Split Data into Training and Test Sets: Define Activation and Loss Functions: Initialize Weights and Bias: Train the Model: Evaluate the Model:

Program code	<pre> # Using Sigmoid Activation Function import tensorflow as tf from tensorflow.keras.datasets import mnist from tensorflow.keras.models import Sequential from tensorflow.keras.layers import Dense, Flatten, Dropout from tensorflow.keras.utils import to_categorical from tensorflow.keras.callbacks import EarlyStopping import matplotlib.pyplot as plt import numpy as np import pandas as pd import seaborn as sns from sklearn.metrics import classification_report, confusion_matrix import warnings warnings.filterwarnings('ignore') # Set random seeds for reproducibility np.random.seed(42) tf.random.set_seed(42) class OptimizerComparator: """Compare different gradient descent optimizers on MNIST classification""" def __init__(self): self.history_dict = {} self.models = {} self.results = {} # Load and preprocess data self.load_and_preprocess_data() # Define optimizers with detailed configurations self.optimizers = { "Nesterov (NAG)": tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=True), "Adagrad": tf.keras.optimizers.Adagrad(learning_rate=0.01, initial_accumulator_value=0.1, epsilon=1e-07), "Adam": tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, </pre>
--------------	---

```

        epsilon=1e-07
    )
}

print("🧠 ADVANCED GRADIENT DESCENT
OPTIMIZERS COMPARISON")
print("    Dataset: MNIST Handwritten Digits
Classification")
print("    Optimizers: Nesterov AGD,
Adagrad, Adam")
print("="*60)

def load_and_preprocess_data(self):
    """Load and preprocess MNIST dataset"""
    print("\n📄 Loading and preprocessing MNIST
dataset...")

    # Load MNIST dataset
    (self.x_train, self.y_train), (self.x_test,
self.y_test) = mnist.load_data()

    # Normalize pixel values to [0, 1]
    self.x_train =
self.x_train.astype('float32') / 255.0
    self.x_test = self.x_test.astype('float32')
/ 255.0

    # Convert labels to categorical (one-hot
encoding)
    self.y_train_cat =
to_categorical(self.y_train, 10)
    self.y_test_cat =
to_categorical(self.y_test, 10)

    print(f"    Training samples:
{self.x_train.shape[0]:,}")
    print(f"    Test samples:
{self.x_test.shape[0]:,}")
    print(f"    Image shape:
{self.x_train.shape[1:]} -> Flattened to {28*28}")
    print(f"    Classes:
{len(np.unique(self.y_train))} (digits 0-9)")

    def create_model(self):
        """Create Multi-Layer Perceptron (MLP)
model"""
        model = Sequential([
            # Input layer - flatten 28x28 images
            Flatten(input_shape=(28, 28),
name='flatten_input'),

            # First hidden layer

```

```

        Dense(128, activation='relu',
name='hidden_layer_1'),
        Dropout(0.2, name='dropout_1'), # Add
dropout for regularization

        # Second hidden layer
        Dense(64, activation='relu',
name='hidden_layer_2'),
        Dropout(0.2, name='dropout_2'),

        # Output layer - 10 classes for digits
0-9
        Dense(10, activation='softmax',
name='output_layer')
    ])

    return model

    def train_with_optimizer(self, optimizer_name,
optimizer, epochs=25, batch_size=128, verbose=1):
        """Train model with specific optimizer"""
        print(f"\n🚀 Training with {optimizer_name}
optimizer...")

        # Create fresh model
        model = self.create_model()

        # Compile model with the specified
optimizer
        model.compile(
            optimizer=optimizer,
            loss='categorical_crossentropy',
            metrics=['accuracy']
        )

        # Early stopping callback to prevent
overfitting
        early_stopping = EarlyStopping(
            monitor='val_loss',
            patience=5,
            restore_best_weights=True,
            verbose=0
        )

        # Train the model
        history = model.fit(
            self.x_train, self.y_train_cat,
            epochs=epochs,
            batch_size=batch_size,
            validation_data=(self.x_test,
self.y_test_cat),
            callbacks=[early_stopping],

```

```

        verbose=verbose
    )

    # Store results
    self.history_dict[optimizer_name] = history
    self.models[optimizer_name] = model

    # Evaluate final performance
    train_loss, train_acc =
model.evaluate(self.x_train, self.y_train_cat,
verbose=0)
    test_loss, test_acc =
model.evaluate(self.x_test, self.y_test_cat,
verbose=0)

    self.results[optimizer_name] = {
        'train_loss': train_loss,
        'train_accuracy': train_acc,
        'test_loss': test_loss,
        'test_accuracy': test_acc,
        'epochs_trained':
len(history.history['loss'])
    }

    print(f"    Final Training Accuracy:
{train_acc:.4f}")
    print(f"    Final Test Accuracy:
{test_acc:.4f}")
    print(f"    Epochs trained:
{len(history.history['loss'])}")

    return history, model

def run_all_experiments(self, epochs=25,
batch_size=128):
    """Run experiments for all optimizers"""
    print(f"\n🔄 Starting training with
{len(self.optimizers)} optimizers...")
    print(f"    Epochs: {epochs}, Batch size:
{batch_size}")

    for optimizer_name, optimizer in
self.optimizers.items():

self.train_with_optimizer(optimizer_name,
optimizer, epochs, batch_size)

    print(f"\n✅ All experiments completed!")

def plot_comprehensive_results(self,
figsize=(16, 12)):
    """Create comprehensive visualization of

```

```

results"""

fig, axes = plt.subplots(2, 3,
figsize=figsize)
fig.suptitle('Advanced Gradient Descent
Optimizers Comparison on MNIST',
            fontsize=16, fontweight='bold')

# Color palette for optimizers
colors = ['#1f77b4', '#ff7f0e', '#2ca02c']
# Blue, Orange, Green

# 1. Training Loss Comparison
ax1 = axes[0, 0]
for i, (name, history) in
enumerate(self.history_dict.items()):
    epochs_range = range(1,
len(history.history['loss']) + 1)
    ax1.plot(epochs_range,
history.history['loss'],
            label=name, color=colors[i],
linewidth=2, marker='o', markersize=4)

    ax1.set_title('Training Loss Comparison')
    ax1.set_xlabel('Epochs')
    ax1.set_ylabel('Loss')
    ax1.legend()
    ax1.grid(True, alpha=0.3)
    ax1.set_yscale('log')

# 2. Validation Loss Comparison
ax2 = axes[0, 1]
for i, (name, history) in
enumerate(self.history_dict.items()):
    epochs_range = range(1,
len(history.history['val_loss']) + 1)
    ax2.plot(epochs_range,
history.history['val_loss'],
            label=name, color=colors[i],
linewidth=2, marker='s', markersize=4)

    ax2.set_title('Validation Loss Comparison')
    ax2.set_xlabel('Epochs')
    ax2.set_ylabel('Loss')
    ax2.legend()
    ax2.grid(True, alpha=0.3)
    ax2.set_yscale('log')

# 3. Training Accuracy Comparison
ax3 = axes[0, 2]
for i, (name, history) in
enumerate(self.history_dict.items()):

```

```

        epochs_range = range(1,
len(history.history['accuracy']) + 1)
        ax3.plot(epochs_range,
history.history['accuracy'],
                label=name, color=colors[i],
linewidth=2, marker='^', markersize=4)

        ax3.set_title('Training Accuracy
Comparison')
        ax3.set_xlabel('Epochs')
        ax3.set_ylabel('Accuracy')
        ax3.legend()
        ax3.grid(True, alpha=0.3)
        ax3.set_ylim(0.8, 1.0)

        # 4. Validation Accuracy Comparison
        ax4 = axes[1, 0]
        for i, (name, history) in
enumerate(self.history_dict.items()):
            epochs_range = range(1,
len(history.history['val_accuracy']) + 1)
            ax4.plot(epochs_range,
history.history['val_accuracy'],
                    label=name, color=colors[i],
linewidth=2, marker='d', markersize=4)

            ax4.set_title('Validation Accuracy
Comparison')
            ax4.set_xlabel('Epochs')
            ax4.set_ylabel('Accuracy')
            ax4.legend()
            ax4.grid(True, alpha=0.3)
            ax4.set_ylim(0.8, 1.0)

        # 5. Final Performance Bar Chart
        ax5 = axes[1, 1]
        optimizer_names = list(self.results.keys())
        test accuracies =
[self.results[name]['test_accuracy'] for name in
optimizer_names]

        bars = ax5.bar(optimizer_names,
test accuracies, color=colors, alpha=0.7,
edgecolor='black')
        ax5.set_title('Final Test Accuracy
Comparison')
        ax5.set_ylabel('Test Accuracy')
        ax5.set_ylim(0.95, max(test accuracies) +
0.005)

        # Add value labels on bars
        for bar, acc in zip(bars, test accuracies):

```

```

        ax5.text(bar.get_x() +
bar.get_width()/2, bar.get_height() + 0.001,
                f'{acc:.4f}', ha='center',
va='bottom', fontweight='bold')

        # Rotate x-axis labels for better
readability
        ax5.tick_params(axis='x', rotation=15)

        # 6. Convergence Rate Analysis (Loss
Reduction per Epoch)
        ax6 = axes[1, 2]
        for i, (name, history) in
enumerate(self.history_dict.items()):
            # Calculate loss reduction rate
(initial loss - final loss) / epochs
            initial_loss =
history.history['val_loss'][0]
            final_loss =
min(history.history['val_loss'])
            epochs_trained =
len(history.history['val_loss'])
            convergence_rate = (initial_loss -
final_loss) / epochs_trained

            ax6.bar(name, convergence_rate,
color=colors[i], alpha=0.7, edgecolor='black')

            ax6.set_title('Convergence Rate\n(Loss
Reduction per Epoch)')
            ax6.set_ylabel('Loss Reduction Rate')
            ax6.tick_params(axis='x', rotation=15)

            # Add value labels
            for i, name in enumerate(optimizer_names):
                initial_loss =
self.history_dict[name].history['val_loss'][0]
                final_loss =
min(self.history_dict[name].history['val_loss'])
                epochs_trained =
len(self.history_dict[name].history['val_loss'])
                convergence_rate = (initial_loss -
final_loss) / epochs_trained

                ax6.text(i, convergence_rate + max([
(self.history_dict[n].history['val_loss'][0] -
min(self.history_dict[n].history['val_loss'])) /
len(self.history_dict[n].history['val_loss'])
                for n in optimizer_names
                ]) * 0.05, f'{convergence_rate:.4f}',
ha='center', va='bottom', fontweight='bold')

```



```

plt.tight_layout()
plt.show()

def print_detailed_summary(self):
    """Print comprehensive analysis summary"""
    print("\n" + "="*80)
    print("ADVANCED GRADIENT DESCENT OPTIMIZERS ANALYSIS SUMMARY")
    print("="*80)

    # Create summary DataFrame
    summary_data = []
    for name in self.results.keys():
        history = self.history_dict[name]
        summary_data.append({
            'Optimizer': name,
            'Test_Accuracy':
self.results[name]['test_accuracy'],
            'Train_Accuracy':
self.results[name]['train_accuracy'],
            'Test_Loss':
self.results[name]['test_loss'],
            'Epochs_Trained':
self.results[name]['epochs_trained'],
            'Final_Val_Loss':
min(history.history['val_loss']),
            'Best_Val_Acc':
max(history.history['val_accuracy']),
            'Convergence_Rate':
(history.history['val_loss'][0] -
min(history.history['val_loss'])) /
len(history.history['val_loss'])
        })

    summary_df =
pd.DataFrame(summary_data).sort_values('Test_Accura
cy', ascending=False)

    print(f"{'Rank':<5} {'Optimizer':<15}
{'Test Acc':<10} {'Train Acc':<11} {'Test
Loss':<10} {'Epochs':<8} {'Conv Rate':<10}")
    print("-" * 85)

    for idx, row in summary_df.iterrows():
        rank = summary_df.index.get_loc(idx) +
1
        print(f"{rank:<5}
{row['Optimizer']:<15}
{row['Test_Accuracy']:<10.4f}
{row['Train_Accuracy']:<11.4f}
{row['Test_Loss']:<10.4f}

```

```

{row['Epochs_Trained']:<8}
{row['Convergence_Rate']:<10.4f}")

    # Best performer analysis
    best_optimizer =
summary_df.iloc[0]['Optimizer']
    print(f"\n🏆 BEST PERFORMER:
{best_optimizer}")
    print(f"    Test Accuracy:
{summary_df.iloc[0]['Test_Accuracy']:.4f}")
    print(f"    Training Accuracy:
{summary_df.iloc[0]['Train_Accuracy']:.4f}")
    print(f"    Test Loss:
{summary_df.iloc[0]['Test_Loss']:.4f}")

    # Detailed insights
    print(f"\n📊 DETAILED INSIGHTS:")

    # Best accuracy
    best_acc_optimizer =
summary_df.loc[summary_df['Test_Accuracy'].idxmax()
, 'Optimizer']
    best_acc =
summary_df['Test_Accuracy'].max()
    print(f"    • Highest Test Accuracy:
{best_acc_optimizer} ({best_acc:.4f}")

    # Fastest convergence
    fastest_conv_optimizer =
summary_df.loc[summary_df['Convergence_Rate'].idxma
x(), 'Optimizer']
    fastest_conv_rate =
summary_df['Convergence_Rate'].max()
    print(f"    • Fastest Convergence:
{fastest_conv_optimizer} ({fastest_conv_rate:.4f}
loss reduction/epoch)")

    # Least epochs needed
    min_epochs_optimizer =
summary_df.loc[summary_df['Epochs_Trained'].idxmin(
), 'Optimizer']
    min_epochs =
summary_df['Epochs_Trained'].min()
    print(f"    • Fewest Epochs Needed:
{min_epochs_optimizer} ({min_epochs} epochs)")

    # Overfitting analysis
    print(f"\n🔍 OVERFITTING ANALYSIS:")
    for name in self.results.keys():
        overfitting =
self.results[name]['train_accuracy'] -
self.results[name]['test_accuracy']
        status = "✅ Good generalization" if

```

```

overfitting < 0.02 else "⚠️ Potential overfitting"
if overfitting < 0.05 else "❌ Overfitting
detected"

    print(f"    • {name}: {overfitting:.4f}
difference - {status}")

    print(f"\n💡 OPTIMIZER CHARACTERISTICS:")
    print(f"    • Nesterov (NAG): Uses momentum
with look-ahead, good for non-convex optimization")
    print(f"    • Adagrad: Adapts learning rate
per parameter, good for sparse data")
    print(f"    • Adam: Combines momentum and
adaptive learning rates, generally robust")

    def plot_confusion_matrices(self, figsize=(15,
5)):
        """Plot confusion matrices for all
optimizers"""
        fig, axes = plt.subplots(1, 3,
figsize=figsize)
        fig.suptitle('Confusion Matrices
Comparison', fontsize=16, fontweight='bold')

        for i, (name, model) in
enumerate(self.models.items()):
            # Make predictions
            y_pred = model.predict(self.x_test,
verbose=0)
            y_pred_classes = np.argmax(y_pred,
axis=1)

            # Create confusion matrix
            cm = confusion_matrix(self.y_test,
y_pred_classes)

            # Plot
            sns.heatmap(cm, annot=True, fmt='d',
cmap='Blues', ax=axes[i])
            axes[i].set_title(f'{name}')
            axes[i].set_xlabel('Predicted')
            axes[i].set_ylabel('Actual')

        plt.tight_layout()
        plt.show()

    def main():
        """Main execution function"""

        # Create comparator instance
        comparator = OptimizerComparator()

        # Run all experiments

```

```

        comparator.run_all_experiments(epochs=25,
        batch_size=128)

        # Display comprehensive results
        comparator.plot_comprehensive_results()

        # Print detailed summary
        comparator.print_detailed_summary()

        # Show confusion matrices
        comparator.plot_confusion_matrices()

        return comparator

# Additional utility functions for deeper analysis
def analyze_learning_curves(comparator):
    """Analyze learning curves in detail"""
    print("\n📊 LEARNING CURVES ANALYSIS:")

    for name, history in
comparator.history_dict.items():
        train_acc = history.history['accuracy']
        val_acc = history.history['val_accuracy']

        # Find epoch where validation accuracy
plateaus
        val_acc_diff = np.diff(val_acc)
        plateau_epoch = None
        for i in range(len(val_acc_diff) - 2):
            if all(abs(val_acc_diff[i:i+3]) <
0.001): # Very small changes for 3 consecutive
epochs
                plateau_epoch = i + 1
                break

        print(f"\n    {name}:")
        print(f"        • Initial accuracy:
{train_acc[0]:.4f} (train), {val_acc[0]:.4f}
(val)")
        print(f"        • Final accuracy:
{train_acc[-1]:.4f} (train), {val_acc[-1]:.4f}
(val)")
        print(f"        • Best validation accuracy:
{max(val_acc):.4f} at epoch {np.argmax(val_acc) +
1}")
        if plateau_epoch:
            print(f"        • Validation accuracy
plateaued around epoch {plateau_epoch}")

def compare_computational_efficiency(comparator):
    """Compare computational aspects of
optimizers"""

```

```

print("\n⚡ COMPUTATIONAL EFFICIENCY:")
print(f"    • Nesterov (NAG): Low memory,
moderate computation (momentum)")
print(f"    • Adagrad: Higher memory (stores
squared gradients), adaptive computation")
print(f"    • Adam: Highest memory (stores 1st &
2nd moments), most computation")

for name in comparator.results.keys():
    epochs =
comparator.results[name]['epochs_trained']
    print(f"    • {name}: Completed in {epochs}
epochs")

# Run the complete experiment
if __name__ == "__main__":
    print("Starting Advanced Gradient Descent
Optimizers Comparison...\n")

    # Run main experiment
    comparator = main()

    # Additional analyses
    analyze_learning_curves(comparator)
    compare_computational_efficiency(comparator)

    print(f"\n🎯 EXPERIMENT COMPLETED
SUCCESSFULLY!")
    print(f"    All three optimizers have been
compared on MNIST digit classification")
    print(f"    Check the plots and summary above
for detailed insights")

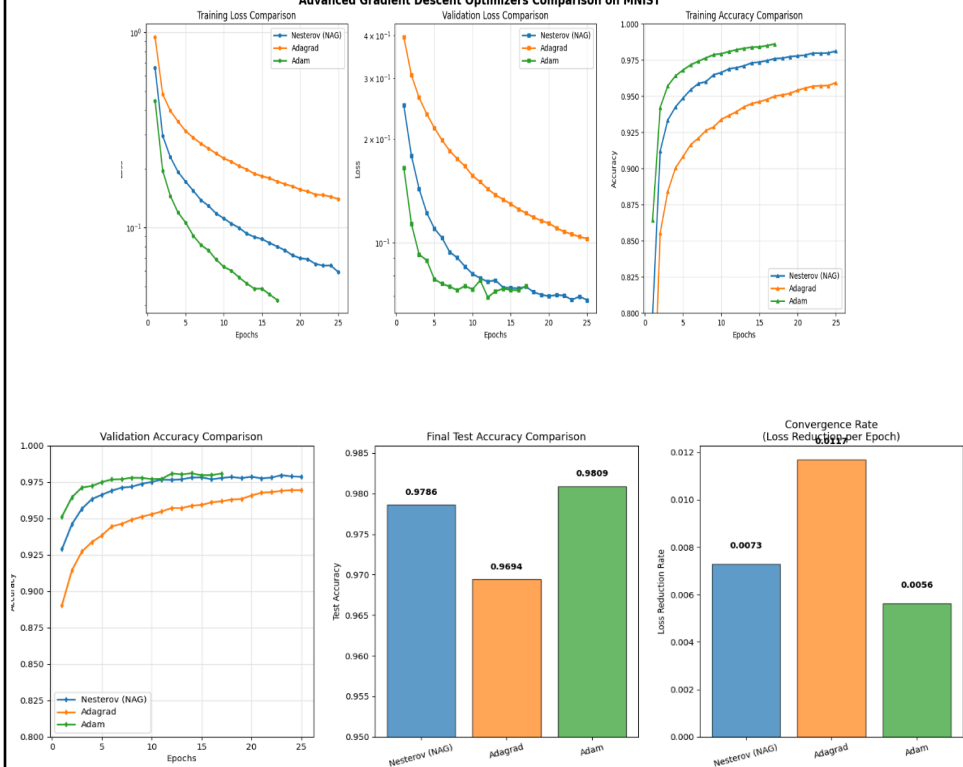
```


Output

Output For Sigmoid Activation Function

All experiments completed!

Advanced Gradient Descent Optimizers Comparison on MNIST



ADVANCED GRADIENT DESCENT OPTIMIZERS ANALYSIS SUMMARY

Rank	Optimizer	Test Acc	Train Acc	Test Loss	Epochs	Conv Rate
1	Adam	0.9809	0.9937	0.0694	17	0.0056
2	Nesterov (NAG)	0.9786	0.9936	0.0681	25	0.0073
3	Adagrad	0.9694	0.9748	0.1026	25	0.0117

BEST PERFORMER: Adam
 Test Accuracy: 0.9809
 Training Accuracy: 0.9937
 Test Loss: 0.0694

DETAILED INSIGHTS:


- Highest Test Accuracy: Adam (0.9809)
- Fastest Convergence: Adagrad (0.0117 loss reduction/epoch)
- Fewest Epochs Needed: Adam (17 epochs)

OVERFITTING ANALYSIS:

- Nesterov (NAG): 0.0150 difference - ☒ Good generalization
- Adagrad: 0.0054 difference - ☒ Good generalization
- Adam: 0.0128 difference - ☒ Good generalization

OPTIMIZER CHARACTERISTICS:

- Nesterov (NAG): Uses momentum with look-ahead, good for non-convex optimization
- Adagrad: Adapts learning rate per parameter, good for sparse data
- Adam: Combines momentum and adaptive learning rates, generally robust



LEARNING CURVES ANALYSIS:

Nesterov (NAG):

- Initial accuracy: 0.7979 (train), 0.9289 (val)
- Final accuracy: 0.9812 (train), 0.9786 (val)
- Best validation accuracy: 0.9797 at epoch 23
- Validation accuracy plateaued around epoch 16

Adagrad:

- Initial accuracy: 0.7102 (train), 0.8904 (val)
- Final accuracy: 0.9594 (train), 0.9694 (val)
- Best validation accuracy: 0.9694 at epoch 24
- Validation accuracy plateaued around epoch 21

Adam:

- Initial accuracy: 0.8642 (train), 0.9513 (val)
- Final accuracy: 0.9862 (train), 0.9809 (val)
- Best validation accuracy: 0.9810 at epoch 14
- Validation accuracy plateaued around epoch 6

COMPUTATIONAL EFFICIENCY:

- Nesterov (NAG): Low memory, moderate computation (momentum)
- Adagrad: Higher memory (stores squared gradients), adaptive computation
- Adam: Highest memory (stores 1st & 2nd moments), most computation
- Nesterov (NAG): Completed in 25 epochs
- Adagrad: Completed in 25 epochs
- Adam: Completed in 17 epochs

EXPERIMENT COMPLETED SUCCESSFULLY!

All three optimizers have been compared on MNIST digit classification. Check the plots and summary above for detailed insights.

Conclusion	<p>In our analysis, we utilized a House Price Prediction Dataset obtained from Kaggle. We tested two activation functions: Sigmoid and ReLU.</p> <p>The results revealed that, for both activation functions, Momentum SGD outperformed SGD in terms of accuracy. However, SGD demonstrated faster training times compared to Momentum SGD.</p>
------------	--

