

Semester	B.E. Semester VII
Subject	Deep Learning
Subject Professor Incharge	Dr. Nayana Mahajan
Laboratory	M201B

Student Name	Harsh Jain	Division	B
Roll Number	22108B0054	Batch	4
Grade and Subject			
Teacher's Signature			

Experiment Number	6
Experiment Title	Image Classification Using Convolutional Neural Network (CNN)
Resources / Apparatus Required	Software: Google Colab, Python

Program code	<pre> # Colab-ready MNIST denoising autoencoder (optimized) # Run in Colab: # 1) Runtime -&gt; Change runtime type -&gt; GPU # 2) (Optional) Mount Drive to save outputs # 3) Run this cell  import os import numpy as np import matplotlib.pyplot as plt import tensorflow as tf from tensorflow.keras import layers, models, callbacks, optimizers, losses from tensorflow.keras.datasets import mnist import datetime  # ----- Config ----- NOISE_FACTOR = 0.4 BATCH_SIZE = 256 # larger batch for GPU EPOCHS = 60 AUTOTUNE = tf.data.AUTOTUNE NUM_DISPLAY = 10 SAVE_DIR = "autoencoder_with_graph_outp ut" USE_MIXED_PRECISION = True # set False if using CPU or older GPU </pre>
--------------	--

```

SEED = 42
MODEL_NAME =
"mnist_denoiser_v2"

os.makedirs(SAVE_DIR,
exist_ok=True)

# -----
Reproducibility
-----

np.random.seed(SEED)
tf.random.set_seed(SEED)

# ----- Mixed
precision -----
if USE_MIXED_PRECISION:
    try:
        # Use the correct
        method for setting the mixed
        precision policy
        policy =
        tf.keras.mixed_precision.Pol
        icy('mixed_float16')

        tf.keras.mixed_precision.set
        _global_policy(policy)
    except AttributeError:
        # Fallback for older
        TF versions if necessary
        print("Could not set
        global mixed precision
        policy. Using
        experimental.")
    try:
        from
        tensorflow.keras.mixed_preci

```

```

sion import experimental as
mixed_precision

    policy =
mixed_precision.Policy('mixed_float16')

mixed_precision.set_policy(policy)

except Exception as
e:

    print(f"Could
not set experimental mixed
precision policy: {e}")

USE_MIXED_PRECISION = False
# Disable if neither works

# ----- Load &
preprocess -----
(x_train, _), (x_test, _) =
mnist.load_data()
x_train =
x_train.astype("float32") /
255.0
x_test =
x_test.astype("float32") /
255.0
x_train =
np.expand_dims(x_train, -1)
# (N, 28, 28, 1)
x_test =
np.expand_dims(x_test, -1)

# Add noise
rng =
np.random.RandomState(SEED)

```

```
x_train_noisy = x_train +
NOISE_FACTOR *
rng.normal(size=x_train.shape)
x_test_noisy = x_test +
NOISE_FACTOR *
rng.normal(size=x_test.shape)
x_train_noisy =
np.clip(x_train_noisy, 0.0,
1.0).astype("float32")
x_test_noisy =
np.clip(x_test_noisy, 0.0,
1.0).astype("float32")

# ----- tf.data
pipeline -----
def make_dataset(clean,
noisy, batch_size,
shuffle=True):
    ds =
tf.data.Dataset.from_tensor_
slices((noisy, clean))
    if shuffle:
        ds =
ds.shuffle(10000, seed=SEED)
    ds =
ds.batch(batch_size,
drop_remainder=False)
    ds = ds.cache()
# keep in memory (MNIST
small)
    ds =
ds.prefetch(AUTOTUNE)
    return ds
```

```

train_ds =
make_dataset(x_train,
x_train_noisy, BATCH_SIZE,
shuffle=True)
val_ds =
make_dataset(x_test,
x_test_noisy, BATCH_SIZE,
shuffle=False)

# ----- Model
builder -----
def
build_autoencoder(input_shape=(28,28,1),
bottleneck_dim=8):
    inp =
layers.Input(shape=input_shape)

    # ENCODER
    x = layers.Conv2D(32, 3,
strides=1,
padding='same')(inp)
    x =
layers.BatchNormalization()(x)
    x =
layers.LeakyReLU(0.2)(x)
    x = layers.Conv2D(64, 3,
strides=2,
padding='same')(x) # 14x14
    x =
layers.BatchNormalization()(x)
    x =
layers.LeakyReLU(0.2)(x)

```

```

        x = layers.Conv2D(128,
3, strides=2,
padding='same')(x) # 7x7
        x =
layers.BatchNormalization()(
x)
        x =
layers.LeakyReLU(0.2)(x)
        x = layers.Conv2D(256,
3, strides=2,
padding='same')(x) # 4x4
(approx)
        x =
layers.BatchNormalization()(
x)
        x =
layers.LeakyReLU(0.2)(x)

        # Bottleneck conv
        encoded =
layers.Conv2D(128, 3,
padding='same')(x) #
compressed representation

        # DECODER
        # Use UpSampling2D and
Conv2D to control spatial
dimensions
        x =
layers.UpSampling2D((2,2))(e
ncoded) # 4x4 -> 8x8
        x = layers.Conv2D(128,
3, padding='same')(x)
        x =
layers.BatchNormalization()(
x)

```

```

x =
layers.LeakyReLU(0.2)(x)

x =
layers.UpSampling2D((2,2))(x)
) # 8x8 -> 16x16
x = layers.Conv2D(64, 3,
padding='same')(x)
x =
layers.BatchNormalization()(
x)
x =
layers.LeakyReLU(0.2)(x)

x =
layers.UpSampling2D((2,2))(x)
) # 16x16 -> 32x32
x = layers.Conv2D(32, 3,
padding='same')(x)
x =
layers.BatchNormalization()(
x)
x =
layers.LeakyReLU(0.2)(x)

# Crop to 28x28
x =
layers.Cropping2D(cropping=(
(2, 2), (2, 2)))(x) # Crop 2
pixels from top, bottom,
left, right

# Final output layer
x = layers.Conv2D(1, 3,
strides=1, padding='same',
activation='sigmoid',

```



```

name='output_sigmoid')(x) #
keep at 28x28

    # If mixed precision is
    active, outputs may be
    float16 -> cast to float32
    for loss computation &
    visualization

    # Only cast if mixed
    precision is actually
    enabled after the checks
    above
    if USE_MIXED_PRECISION:
        x =
layers.Lambda(lambda t:
tf.cast(t, tf.float32),
name='out_cast')(x)

    model =
models.Model(inp, x,
name="denoising_autoencoder"
)

    return model

autoencoder =
build_autoencoder()
autoencoder.summary()

# ----- Optimizer
& Loss -----
base_lr = 1e-3
if USE_MIXED_PRECISION:
    opt =
optimizers.Adam(learning_rat
e=base_lr)

```

```

        # Keras handles dynamic
loss scaling automatically
with mixed precision policy
in TF 2.10+
else:
    opt =
optimizers.Adam(learning_lea
rning_rate=base_lr)

# For denoising, MSE often
gives good reconstruction;
you can switch to
binary_crossentropy if
desired
loss_fn =
losses.MeanSquaredError()

autoencoder.compile(optimiz
r=opt, loss=loss_fn,
metrics=['mae'])

# ----- Callbacks
-----

timestamp =
datetime.datetime.now().strf
time("%Y%m%d-%H%M%S")
tb_logdir =
os.path.join(SAVE_DIR,
"tb_logs", MODEL_NAME + "_"
+ timestamp)
os.makedirs(tb_logdir,
exist_ok=True)

cb_list = [

callbacks.ModelCheckpoint(

```

```

filepath=os.path.join(SAVE_D
IR, MODEL_NAME +
"_best.h5"),
        monitor='val_loss',
save_best_only=True,
verbose=1
    ),

callbacks.EarlyStopping(moni
tor='val_loss', patience=8,
restore_best_weights=True,
verbose=1),

callbacks.ReduceLROnPlateau(
monitor='val_loss',
factor=0.5, patience=4,
min_lr=1e-6, verbose=1),

callbacks.TensorBoard(log_di
r=tb_logdir,
histogram_freq=1,
write_images=True)
]

# ----- Training
-----

history = autoencoder.fit(
    train_ds,
    validation_data=val_ds,
    epochs=EPOCHS,
    callbacks=cb_list,
    verbose=2
)

# ----- Save
final model -----

```

```
# Add .keras extension for
SavedModel format
saved_model_filepath =
os.path.join(SAVE_DIR,
MODEL_NAME +
"_savedmodel.keras")
autoencoder.save(saved_model
_filepath,
include_optimizer=False)
autoencoder.save(os.path.joi
n(SAVE_DIR, MODEL_NAME +
".h5"))

print("Saved model to:",
saved_model_filepath)

# ----- Plot loss
curves -----
plt.figure(figsize=(8,5))
plt.plot(history.history['lo
ss'], label='train_loss')
plt.plot(history.history['va
l_loss'], label='val_loss')
plt.xlabel('Epoch');
plt.ylabel('Loss');
plt.grid(True); plt.legend()
plt.title('Training &
Validation Loss')
plt.tight_layout()
plt.savefig(os.path.join(SAV
E_DIR, "loss_curve.png"))
plt.show()

# ----- Denoise
and visualize
-----
```

```

num_display =
min(NUM_DISPLAY,
x_test.shape[0])
decoded =
autoencoder.predict(x_test_noisy[:num_display])

def
display_images_grid(orig,
noisy, denoised,
n=num_display,
savepath=None):

plt.figure(figsize=(16,6))
    for i in range(n):
        # original
        ax = plt.subplot(3,
n, i+1)

plt.imshow(orig[i].squeeze()
, cmap='gray');
plt.axis('off')
        if i == 0:
plt.title("Original")
        # noisy
        ax = plt.subplot(3,
n, i+1+n)

plt.imshow(noisy[i].squeeze()
), cmap='gray');
plt.axis('off')
        if i == 0:
plt.title("Noisy")
        # denoised
        ax = plt.subplot(3,
n, i+1+2*n)

```

```
plt.imshow(denoised[i].squeeze(), cmap='gray');
plt.axis('off')
    if i == 0:
plt.title("Denoised")
    plt.tight_layout()
    if savepath:

plt.savefig(savepath)
    plt.show()

vis_path =
os.path.join(SAVE_DIR,
"denoising_results.png")
display_images_grid(x_test,
x_test_noisy, decoded,
n=num_display,
savepath=vis_path)

print("Saved visuals to:",
vis_path)
print("TensorBoard logs
in:", tb_logdir)

# ----- Optional:
write a few images to
TensorBoard for visual
inspection -----
try:
    file_writer =
tf.summary.create_file_writ
er(tb_logdir + "/images")
    with
file_writer.as_default():
        # log first batch of
images
```

```

        for i in
range(min(3, num_display)):


tf.summary.image(f"orig_{i}"
, np.expand_dims(x_test[i],
axis=0), step=0)

tf.summary.image(f"noisy_{i}"
",
np.expand_dims(x_test_noisy[
i], axis=0), step=0)

tf.summary.image(f"denoised_
{i}",
np.expand_dims(decoded[i],
axis=0), step=0)
    print("Wrote image
summaries to TensorBoard.")
except Exception as e:
    print("Could not write
TensorBoard images:", e)

```



Original

Noisy

Denoised

Original

Noisy

Denoised

Saved visuals to:  
autoencoder\_with\_graph\_output/denoising\_results.png  
TensorBoard logs in:  
autoencoder\_with\_graph\_output/tb\_logs/mnist\_denoiser\_v2\_2  
0251003-064832

	Wrote image summaries to TensorBoard.																							
	Model: "denoising_autoencoder"																							
	<table><tr><th>Layer (type)</th><th>Output Shape</th></tr><tr><th>Param #</th><th></th></tr><tr><td>input_layer_8 (InputLayer)</td><td>(None, 28, 28, 1)</td></tr><tr><td>conv2d_49 (Conv2D)</td><td>(None, 28, 28, 32)</td></tr><tr><td>batch_normalization_40</td><td>(None, 28, 28, 32)</td></tr><tr><td>leaky_re_lu_40 (LeakyReLU)</td><td>(None, 28, 28, 32)</td></tr><tr><td>conv2d_50 (Conv2D)</td><td>(None, 14, 14, 64)</td></tr><tr><td>batch_normalization_41</td><td>(None, 14, 14, 64)</td></tr><tr><td>leaky_re_lu_41 (LeakyReLU)</td><td>(None, 14, 14, 64)</td></tr><tr><td>conv2d_51 (Conv2D)</td><td>(None, 7, 7, 128)</td></tr><tr><td>batch_normalization_42</td><td>(None, 7, 7, 128)</td></tr></table>		Layer (type)	Output Shape	Param #		input_layer_8 (InputLayer)	(None, 28, 28, 1)	conv2d_49 (Conv2D)	(None, 28, 28, 32)	batch_normalization_40	(None, 28, 28, 32)	leaky_re_lu_40 (LeakyReLU)	(None, 28, 28, 32)	conv2d_50 (Conv2D)	(None, 14, 14, 64)	batch_normalization_41	(None, 14, 14, 64)	leaky_re_lu_41 (LeakyReLU)	(None, 14, 14, 64)	conv2d_51 (Conv2D)	(None, 7, 7, 128)	batch_normalization_42	(None, 7, 7, 128)
Layer (type)	Output Shape																							
Param #																								
input_layer_8 (InputLayer)	(None, 28, 28, 1)																							
conv2d_49 (Conv2D)	(None, 28, 28, 32)																							
batch_normalization_40	(None, 28, 28, 32)																							
leaky_re_lu_40 (LeakyReLU)	(None, 28, 28, 32)																							
conv2d_50 (Conv2D)	(None, 14, 14, 64)																							
batch_normalization_41	(None, 14, 14, 64)																							
leaky_re_lu_41 (LeakyReLU)	(None, 14, 14, 64)																							
conv2d_51 (Conv2D)	(None, 7, 7, 128)																							
batch_normalization_42	(None, 7, 7, 128)																							

	(BatchNormalization)	
	leaky_re_lu_42 (LeakyReLU)	(None, 7, 7, 128)
	0	
	conv2d_52 (Conv2D)	(None, 4, 4, 256)
	295,168	
	batch_normalization_43	(None, 4, 4, 256)
	1,024	
	(BatchNormalization)	
	leaky_re_lu_43 (LeakyReLU)	(None, 4, 4, 256)
	0	
	conv2d_53 (Conv2D)	(None, 4, 4, 128)
	295,040	
	up_sampling2d_11 (UpSampling2D)	(None, 8, 8, 128)
	0	
	conv2d_54 (Conv2D)	(None, 8, 8, 128)
	147,584	
	batch_normalization_44	(None, 8, 8, 128)
	512	
	(BatchNormalization)	
	leaky_re_lu_44 (LeakyReLU)	(None, 8, 8, 128)
	0	
	up_sampling2d_12 (UpSampling2D)	(None, 16, 16, 128)
	0	
	conv2d_55 (Conv2D)	(None, 16, 16, 64)
	73,792	

batch_normalization_45 256 (BatchNormalization)	(None, 16, 16, 64)
leaky_re_lu_45 0 (LeakyReLU)	(None, 16, 16, 64)
up_sampling2d_13 0 (UpSampling2D)	(None, 32, 32, 64)
conv2d_56 18,464 (Conv2D)	(None, 32, 32, 32)
batch_normalization_46 128 (BatchNormalization)	(None, 32, 32, 32)
leaky_re_lu_46 0 (LeakyReLU)	(None, 32, 32, 32)
cropping2d_2 0 (Cropping2D)	(None, 28, 28, 32)
output_sigmoid 289 (Conv2D)	(None, 28, 28, 1)
out_cast 0 (Lambda)	(None, 28, 28, 1)
<p><b>Total params:</b> 925,825 (3.53 MB)</p> <p><b>Trainable params:</b> 924,417 (3.53 MB)</p> <p><b>Non-trainable params:</b> 1,408 (5.50 KB)</p>	

Conclusion	<p>This comprehensive experiment successfully demonstrated that Convolutional Neural Networks (CNNs) significantly outperform traditional dense networks for RGB image classification on the CIFAR-10 dataset, achieving higher accuracy through spatial feature extraction enabled by ReLU activations and multi-class probability distributions provided by Softmax activation. The CNN's ability to preserve spatial relationships, extract hierarchical features across color channels, and maintain translation invariance makes it fundamentally superior to dense networks that treat pixels as independent features, validating the importance of architecture choice for computer vision tasks. The explicit implementation of ReLU for feature extraction and Softmax for classification showcased how these activation functions enable deep networks to learn complex patterns from RGB data while providing interpretable probability outputs for multi-class classification.</p>
------------	--